

Capstone_MovieLens

Puja

01/13/2022

INTRODUCTION/OVERVIEW/EXECUTIVE SUMMARY

This is a data analysis report prompted by the 9th and final section, Capstone, of the edx program: HarvardX-DataScience. The aim of this project is to test its students' overall capabilities by allowing them to apply all the skills they've learnt thus far in an intriguing project surrounding the world of movies! This project expects its students to take a deep dive into the MovieLens data set and familiarize themselves with different components and the different relationships shared between these MovieLens components. They will then have to implement machine learning skills and construct a body of code (algorithm) using a training set in order to try and predict movie ratings. The goal of this project is to ultimately apply this constructed code onto a test set to see whether it is able to correctly predict movie ratings. We will do this by observing the Root Mean Square Error (RMSE) result.

GENERATE GIVEN DATA

The following code shown below is provided by the "HarvardEDX:Data Science - Capstone" course to its students. This body of code allows you to download the "MovieLens" data, as well as create the training and validation sets. Students are required to create a algorithm using the given 'edx' data set and apply it to the given 'validation' data set.

```
#####  
# Create edx set, validation set (final hold-out test set)  
#####  
  
# Note: this process could take a couple of minutes  
  
if(!require(tidyverse)) install.packages("tidyverse", repos = "http://cran.us.r-project.org")  
  
## Loading required package: tidyverse  
  
## -- Attaching packages ----- tidyverse 1.3.1 --  
  
## v ggplot2 3.3.5      v purrr  0.3.4  
## v tibble  3.1.6      v dplyr  1.0.7  
## v tidyr   1.1.4      v stringr 1.4.0  
## v readr   2.1.1      v forcats 0.5.1  
  
## -- Conflicts ----- tidyverse_conflicts() --  
## x dplyr::filter() masks stats::filter()  
## x dplyr::lag()    masks stats::lag()  
  
if(!require(caret)) install.packages("caret", repos = "http://cran.us.r-project.org")  
  
## Loading required package: caret
```

```

## Loading required package: lattice

##
## Attaching package: 'caret'

## The following object is masked from 'package:purrr':
##
## lift

if(!require(data.table)) install.packages("data.table", repos = "http://cran.us.r-project.org")

## Loading required package: data.table

##
## Attaching package: 'data.table'

## The following objects are masked from 'package:dplyr':
##
## between, first, last

## The following object is masked from 'package:purrr':
##
## transpose

library(tidyverse)
library(caret)
library(data.table)

# MovieLens 10M dataset:
# https://grouplens.org/datasets/movielens/10m/
# http://files.grouplens.org/datasets/movielens/ml-10m.zip

dl <- tempfile()
download.file("http://files.grouplens.org/datasets/movielens/ml-10m.zip", dl)

ratings <- fread(text = gsub(":", "\t", readLines(unzip(dl, "ml-10M100K/ratings.dat"))),
  col.names = c("userId", "movieId", "rating", "timestamp"))

movies <- str_split_fixed(readLines(unzip(dl, "ml-10M100K/movies.dat")), "\\:", 3)
colnames(movies) <- c("movieId", "title", "genres")

# if using R 4.0 or later:
movies <- as.data.frame(movies) %>% mutate(movieId = as.numeric(movieId),
  title = as.character(title),
  genres = as.character(genres))

movielens <- left_join(ratings, movies, by = "movieId")

# Validation set will be 10% of MovieLens data
set.seed(1, sample.kind="Rounding") # if using R 3.5 or earlier, use `set.seed(1)`

## Warning in set.seed(1, sample.kind = "Rounding"): non-uniform 'Rounding' sampler
## used

test_index <- createDataPartition(y = movielens$rating, times = 1, p = 0.1, list = FALSE)
edx <- movielens[-test_index,]
temp <- movielens[test_index,]

```

```

# Make sure userId and movieId in validation set are also in edx set
validation <- temp %>%
  semi_join(edx, by = "movieId") %>%
  semi_join(edx, by = "userId")

# Add rows removed from validation set back into edx set
removed <- anti_join(temp, validation)

## Joining, by = c("userId", "movieId", "rating", "timestamp", "title", "genres")
edx <- rbind(edx, removed)

rm(dl, ratings, movies, test_index, temp, movielens, removed)

```

METHODS/ANALYSIS

DATA SURVEILLANCE

I shall now conduct a thorough Data Surveillance of the given data set in order to familiarize myself with it. My observations will be divided into 2 parts; namely: **Exploration** and **Visualization**. As I carry out my surveillance, I shall also attempt to apply data **cleaning** methods where it may see fit to do so.

EXPLORATION

The purpose of this section is to broaden our understanding of the data. This will be done by carrying out numerous data exploration techniques used throughout the course. It will also include the code used in order to answer the quiz “MovieLens_Dataset”, as I found it to be extremely helpful in familiarizing myself to the contents of the data.

Overall summary of edx dataset.

```
summary(edx)
```

```

##      userId      movieId      rating      timestamp
##  Min.   :    1   Min.   :    1   Min.   :0.500   Min.   :7.897e+08
## 1st Qu.:18124   1st Qu.:   648   1st Qu.:3.000   1st Qu.:9.468e+08
## Median :35738   Median :  1834   Median :4.000   Median :1.035e+09
## Mean   :35870   Mean   :   4122   Mean   :3.512   Mean   :1.033e+09
## 3rd Qu.:53607   3rd Qu.:  3626   3rd Qu.:4.000   3rd Qu.:1.127e+09
## Max.   :71567   Max.   : 65133   Max.   :5.000   Max.   :1.231e+09
##      title      genres
## Length:9000055   Length:9000055
## Class :character Class :character
## Mode  :character Mode  :character
##
##
##

```

```
glimpse(edx)
```

```

## Rows: 9,000,055
## Columns: 6
## $ userId      <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, ~

```

```
## $ movieId    <dbl> 122, 185, 292, 316, 329, 355, 356, 362, 364, 370, 377, 420, ~
## $ rating     <dbl> 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, ~
## $ timestamp  <int> 838985046, 838983525, 838983421, 838983392, 838983392, 83898~
## $ title      <chr> "Boomerang (1992)", "Net, The (1995)", "Outbreak (1995)", "S~
## $ genres     <chr> "Comedy|Romance", "Action|Crime|Thriller", "Action|Drama|Sci~
```

Number of different users, movies, ratings, timestamps, titles and genres in edx data set.

```
Distinct_Numbers <- edx %>%
  summarize(D_userId = n_distinct(edx$userId),
            D_movieId = n_distinct(edx$movieId),
            D_rating = n_distinct(edx$rating),
            D_timestamp = n_distinct(edx$timestamp),
            D_title = n_distinct(edx$title),
            D_genres = n_distinct(edx$genres))
Distinct_Numbers
```

```
##   D_userId D_movieId D_rating D_timestamp D_title D_genres
## 1    69878    10677     10    6519590  10676     797
```

Number of movie ratings made for each of the genres in the edx data set.

```
Genre_Ratings <- edx %>%
  summarize(Drama_Ratings = sapply("Drama", function(d) {
    sum(str_detect(edx$genres, d))}),
            Comedy_Ratings = sapply("Comedy", function(c) {
    sum(str_detect(edx$genres, c))}),
            Thriller_Ratings = sapply("Thriller", function(t) {
    sum(str_detect(edx$genres, t))}),
            Romance_Ratings = sapply("Romance", function(r) {
    sum(str_detect(edx$genres, r))})
  )
Genre_Ratings
```

```
##   Drama_Ratings Comedy_Ratings Thriller_Ratings Romance_Ratings
## 1      3910127      3540930      2325899      1712100
```

Movies with the most amount of cumulative ratings

```
Movie_Ratings <- edx %>%
  group_by(title, genres) %>%
  summarize(cumulative_ratings = n()) %>%
  arrange(desc(cumulative_ratings))
```

`summarise()` has grouped output by 'title'. You can override using the `.groups` argument.

```
Movie_Ratings

## # A tibble: 10,677 x 3
## # Groups:   title [10,676]
##   title                                genres                cumulative_ratin~
##   <chr>                                <chr>                <int>
## 1 Pulp Fiction (1994)                  Comedy|Crime|Drama    31362
## 2 Forrest Gump (1994)                  Comedy|Drama|Romance~ 31079
## 3 Silence of the Lambs, The (1991)     Crime|Horror|Thriller 30382
## 4 Jurassic Park (1993)                  Action|Adventure|Sci~ 29360
## 5 Shawshank Redemption, The (1994)     Drama                28015
## 6 Braveheart (1995)                     Action|Drama|War      26212
```

```
## 7 Fugitive, The (1993) Thriller 25998
## 8 Terminator 2: Judgment Day (1991) Action|Sci-Fi 25984
## 9 Star Wars: Episode IV - A New Hope (~ Action|Adventure|Sci~ 25672
## 10 Apollo 13 (1995) Adventure|Drama 24284
## # ... with 10,667 more rows
```

The most popular ratings used

```
Most_Used_Ratings <- edx %>%
  group_by(rating) %>%
  summarize(times_used = n()) %>%
  top_n(10) %>%
  arrange(desc(times_used))
```

```
## Selecting by times_used
```

```
Most_Used_Ratings
```

```
## # A tibble: 10 x 2
##   rating times_used
##   <dbl>     <int>
## 1     4     2588430
## 2     3     2121240
## 3     5     1390114
## 4   3.5      791624
## 5     2      711422
## 6   4.5      526736
## 7     1      345679
## 8   2.5      333010
## 9   1.5      106426
## 10    0.5       85374
```

Genres with most amount of cumulative ratings

```
Genre_Ratings <- edx %>%
  group_by(genres) %>%
  summarize(cumulative_ratings = n()) %>%
  arrange(desc(cumulative_ratings))
Genre_Ratings
```

```
## # A tibble: 797 x 2
##   genres cumulative_ratings
##   <chr>             <int>
## 1 Drama             733296
## 2 Comedy            700889
## 3 Comedy|Romance    365468
## 4 Comedy|Drama      323637
## 5 Comedy|Drama|Romance 261425
## 6 Drama|Romance      259355
## 7 Action|Adventure|Sci-Fi 219938
## 8 Action|Adventure|Thriller 149091
## 9 Drama|Thriller     145373
## 10 Crime|Drama       137387
## # ... with 787 more rows
```

Extracting the year from the edx column 'title' in order to have a separate column called 'titleyear'.

```
library(stringr)
year_from_title <- '\\d{4}(?=\\))'
titleyear = str_extract(edx$title, year_from_title)
```

New edx dataframe with added column 'titleyear'.

```
new_edx <- edx %>%
  add_column(titleyear)
new_edx
```

```
##      userId movieId rating timestamp title
##      1:      1      122      5.0 838985046 Boomerang (1992)
##      2:      1      185      5.0 838983525 Net, The (1995)
##      3:      1      292      5.0 838983421 Outbreak (1995)
##      4:      1      316      5.0 838983392 Stargate (1994)
##      5:      1      329      5.0 838983392 Star Trek: Generations (1994)
##      ---
## 9000051: 32620 33140 3.5 1173562747 Down and Derby (2005)
## 9000052: 40976 61913 3.0 1227767528 Africa addio (1966)
## 9000053: 59269 63141 2.0 1226443318 Rockin' in the Rockies (1945)
## 9000054: 60713 4820 2.0 1119156754 Won't Anybody Listen? (2000)
## 9000055: 64621 39429 2.5 1201248182 Confess (2005)
##      genres titleyear
##      1:      Comedy|Romance 1992
##      2:      Action|Crime|Thriller 1995
##      3: Action|Drama|Sci-Fi|Thriller 1995
##      4:      Action|Adventure|Sci-Fi 1994
##      5: Action|Adventure|Drama|Sci-Fi 1994
##      ---
## 9000051:      Children|Comedy 2005
## 9000052:      Documentary 1966
## 9000053: Comedy|Musical|Western 1945
## 9000054:      Documentary 2000
## 9000055:      Drama|Thriller 2005
```

```
head(new_edx)
```

```
##      userId movieId rating timestamp title
## 1:      1      122      5 838985046 Boomerang (1992)
## 2:      1      185      5 838983525 Net, The (1995)
## 3:      1      292      5 838983421 Outbreak (1995)
## 4:      1      316      5 838983392 Stargate (1994)
## 5:      1      329      5 838983392 Star Trek: Generations (1994)
## 6:      1      355      5 838984474 Flintstones, The (1994)
##      genres titleyear
## 1:      Comedy|Romance 1992
## 2:      Action|Crime|Thriller 1995
## 3: Action|Drama|Sci-Fi|Thriller 1995
## 4:      Action|Adventure|Sci-Fi 1994
## 5: Action|Adventure|Drama|Sci-Fi 1994
## 6:      Children|Comedy|Fantasy 1994
```

Count how many times each year appears

```
new_edx %>% count(titleyear)
```

##	titleyear	n
## 1:	1915	180
## 2:	1916	84
## 3:	1917	32
## 4:	1918	73
## 5:	1919	158
## 6:	1920	575
## 7:	1921	406
## 8:	1922	1825
## 9:	1923	316
## 10:	1924	457
## 11:	1925	2654
## 12:	1926	383
## 13:	1927	4133
## 14:	1928	1336
## 15:	1929	573
## 16:	1930	2472
## 17:	1931	7669
## 18:	1932	3263
## 19:	1933	7675
## 20:	1934	5954
## 21:	1935	6234
## 22:	1936	4120
## 23:	1937	13456
## 24:	1938	7775
## 25:	1939	27422
## 26:	1940	26708
## 27:	1941	23883
## 28:	1942	20051
## 29:	1943	3563
## 30:	1944	11918
## 31:	1945	5838
## 32:	1946	16882
## 33:	1947	6558
## 34:	1948	9878
## 35:	1949	7862
## 36:	1950	17377
## 37:	1951	21789
## 38:	1952	11540
## 39:	1953	18740
## 40:	1954	30089
## 41:	1955	21673
## 42:	1956	16053
## 43:	1957	24557
## 44:	1958	23336
## 45:	1959	30843
## 46:	1960	29637
## 47:	1961	26004
## 48:	1962	33622
## 49:	1963	31489
## 50:	1964	40477
## 51:	1965	23475
## 52:	1966	18284
## 53:	1967	39745

```
## 54:      1968  48313
## 55:      1969  25702
## 56:      1970  27660
## 57:      1971  59092
## 58:      1972  39512
## 59:      1973  50625
## 60:      1974  47460
## 61:      1975  67996
## 62:      1976  45298
## 63:      1977  60035
## 64:      1978  55020
## 65:      1979  84452
## 66:      1980 104842
## 67:      1981  96947
## 68:      1982 121258
## 69:      1983  94372
## 70:      1984 163168
## 71:      1985 139778
## 72:      1986 175548
## 73:      1987 171738
## 74:      1988 171628
## 75:      1989 228426
## 76:      1990 230409
## 77:      1991 196681
## 78:      1992 236834
## 79:      1993 481184
## 80:      1994 671298
## 81:      1995 786762
## 82:      1996 593518
## 83:      1997 429751
## 84:      1998 402187
## 85:      1999 489537
## 86:      2000 382823
## 87:      2001 305705
## 88:      2002 272180
## 89:      2003 211397
## 90:      2004 204811
## 91:      2005 128613
## 92:      2006 103870
## 93:      2007  75788
## 94:      2008  26741
##      titleyear      n
```

VISUALIZATION


In this section, we will build up on what we've learnt thus far using visual aids. I personally find data visualization to be extremely accommodating when it comes to working with great amounts of data as they help in seeing the bigger picture more clearly (pun intended).

Bar Graph of Amount of Movies Each Year

```
Amount_of_Movies_Each_Year <- new_edx %>%
  ggplot(aes(titleyear)) +
  geom_bar(color = "chocolate4", fill = "darkorange1") +
```

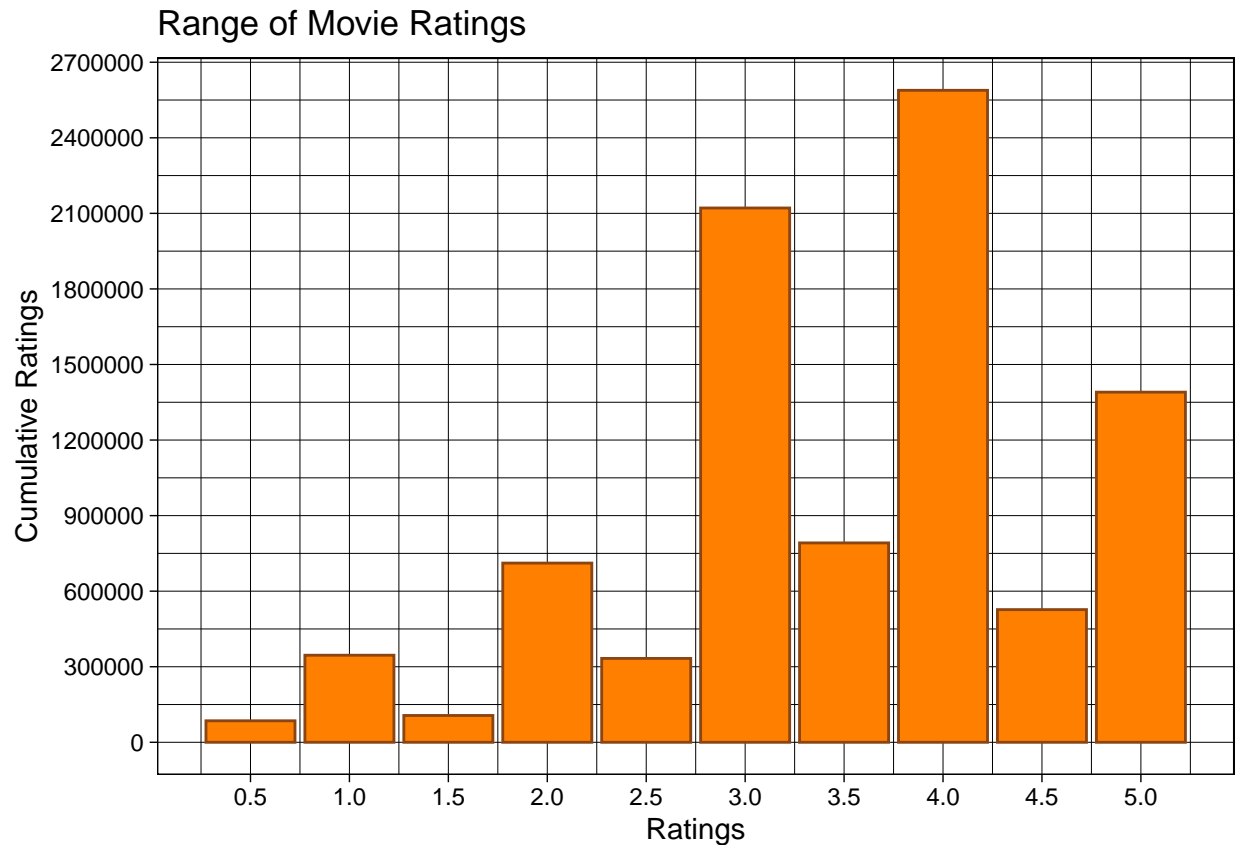


```
labs(x = "Year", y = "Amount of Movies") +
ggtitle("Amount of Movies Each Year") +
theme_linedraw()
Amount_of_Movies_Each_Year
```



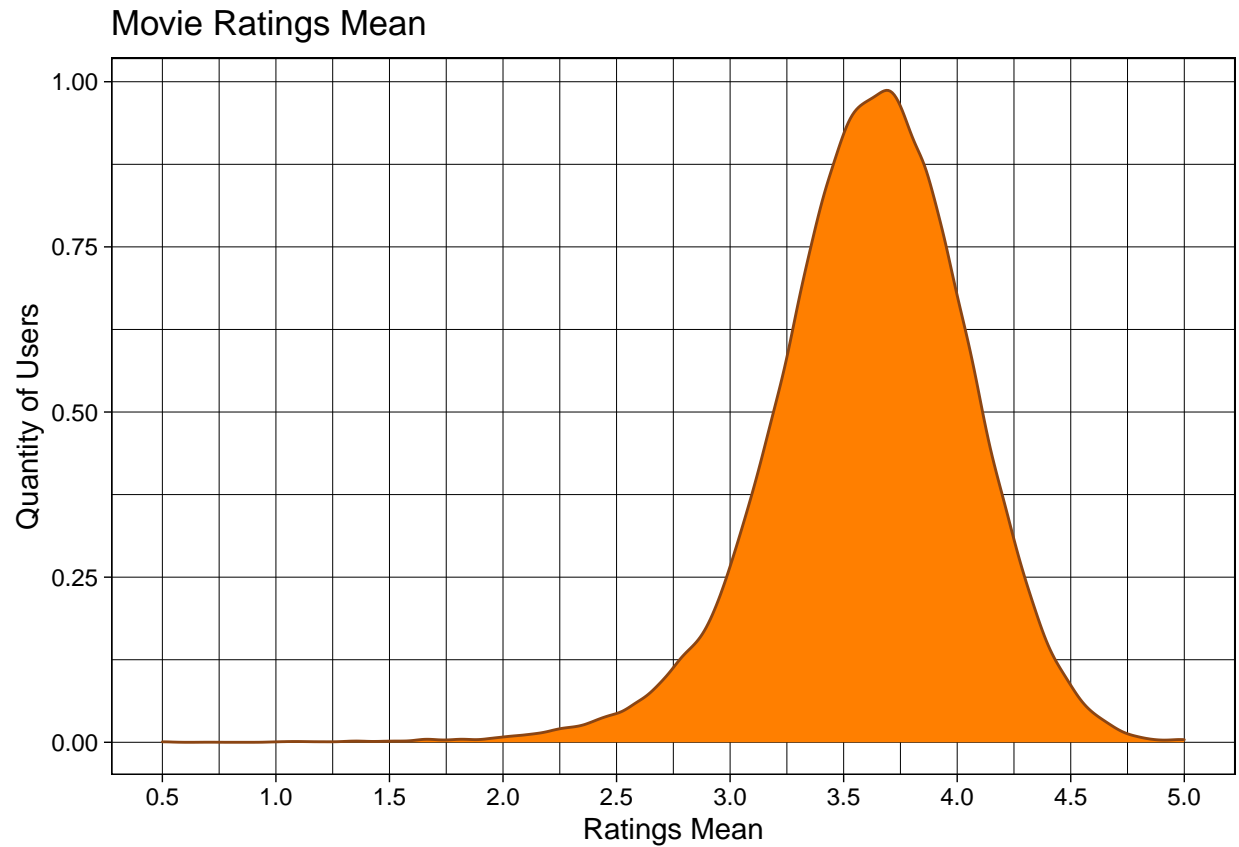
Bar Graph of Range of Movie Ratings

```
Range_of_Movie_Ratings <- edx %>%
  ggplot(aes(rating)) +
  geom_bar(color = "chocolate4", fill = "darkorange1") +
  scale_x_continuous(breaks = c(0.5, 1, 1.5, 2, 2.5, 3, 3.5, 4, 4.5, 5)) +
  scale_y_continuous(breaks = c(0, 300000, 600000, 900000, 1200000, 1500000,
                                1800000, 2100000, 2400000, 2700000)) +
  labs(x = "Ratings", y = "Cumulative Ratings") +
  ggtitle("Range of Movie Ratings") +
  theme_linedraw()
Range_of_Movie_Ratings
```



Density Graph of the Movie Ratings' Mean

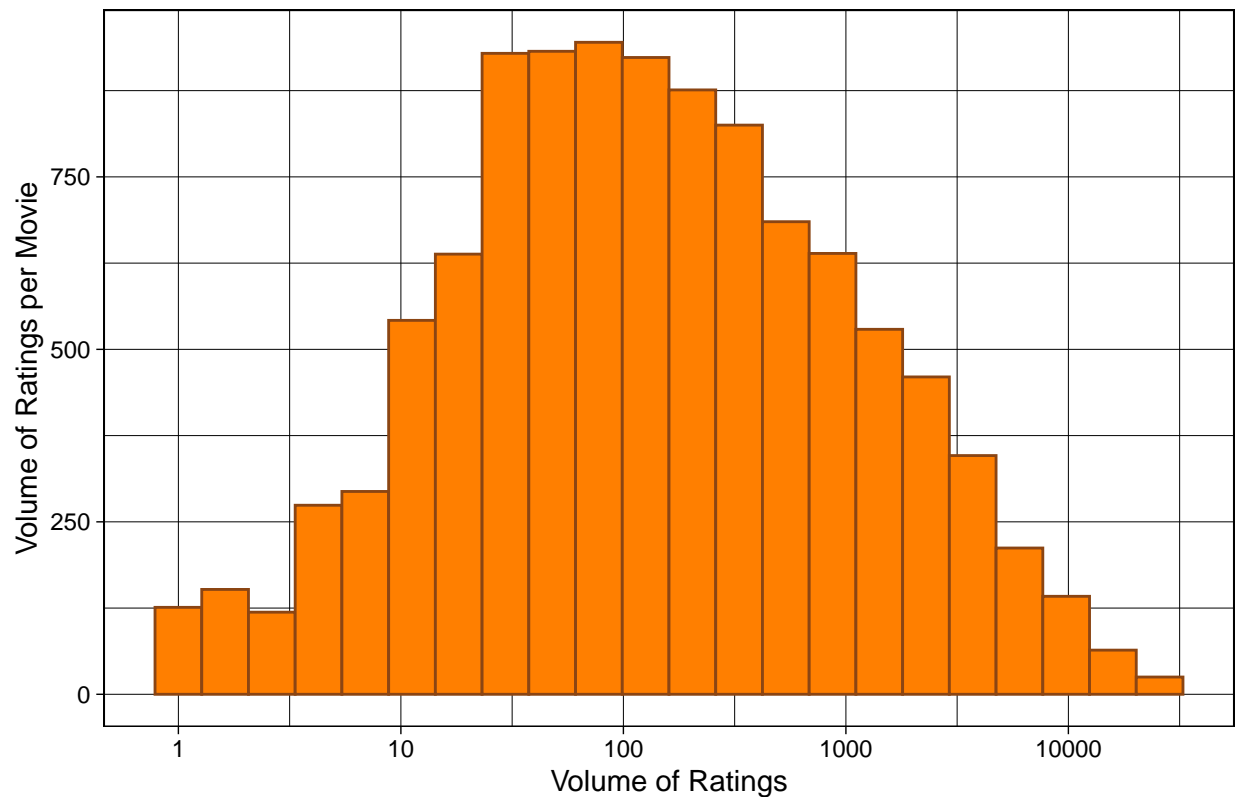
```
Movie_Ratings_Mean <- edx %>%
  group_by(userId) %>%
  summarize(average = mean(rating)) %>%
  ggplot(aes(average)) +
  geom_density(color = "chocolate4", fill = "darkorange1") +
  ggtitle("Movie Ratings Mean") +
  scale_x_continuous(breaks = c(seq(0.5, 5, 0.5))) +
  labs(x = "Ratings Mean", y = "Quantity of Users") +
  theme_linedraw()
Movie_Ratings_Mean
```



Histogram for Ratings Volume per Movie


```
Ratings_Volume_per_Movie <- edx %>%
  count(movieId) %>%
  ggplot(aes(n)) +
  geom_histogram(binwidth = 0.21, color = "chocolate4", fill = "darkorange1") +
  scale_x_log10() +
  scale_y_continuous() +
  labs(x = "Volume of Ratings", y = "Volume of Ratings per Movie") +
  ggtitle("Ratings Volume per Movie") +
  theme_linedraw()
Ratings_Volume_per_Movie
```

Ratings Volume per Movie



Histogram of the Amount of User Ratings

```
Amount_of_User_Ratings <- edx %>%
  group_by(userId) %>%
  summarize(count = n()) %>%
  ggplot(aes(count)) +
  geom_histogram(bins = 21, color = "chocolate4", fill = "darkorange1") +
  scale_x_log10() +
  scale_y_log10() +
  labs(x = "Amount of Ratings", y = "Amount of Users") +
  ggtitle("Amount of User Ratings") +
  theme_linedraw()
Amount_of_User_Ratings
```



MODELING APPROACHES AND INSIGHTS GAINED

We shall now start with the process of modelling our algorithms. Thanks to the various methods of **Data Surveillance** above, we now have an extremely clear idea of the data set that we're working with and it's content. We shall now use the knowledge gained to help construct the algorithm. The accuracy of our work will be determined by the utilization of the Residual Mean Squared Error (RMSE) Formula.

The RMSE formula helps us to calculate the difference (i.e. error) between the predicted and the observed values. Our aim is to obtain an error value of less than 0.8649. We shall apply this RMSE formula to the models below in order to generate their error quantities and see whether we are able to get it below 0.8649.

RMSE FORMULA

```
RMSE_Formula <- function(observed, predicted){
  sqrt(mean((observed - predicted)^2))}
```

I shall now represent "validation\$rating" as "a" in order to save myself from typing it out whenever I have to calculate the RMSE.

```
a <- validation$rating
```

1ST MODEL

The first model which I will be applying to the data set will be the most straightforward application model. It will be implemented by using the mean/average of all the recorded movie ratings in the data set.

Compute the mean/average of all recorded ratings belonging to the edx data set.

```
mu_cap <- mean(edx$rating)
```

RMSE of Model 1

```
RMSE_MODEL_1 <- RMSE_Formula(a, mu_cap)
RMSE_MODEL_1
```

```
## [1] 1.061202
```

INSIGHTS GAINED FROM MODEL 1:

As predicted, our RMSE result for Model 1 is not nearly as close to our desired reading. This may be because the ratings alone is not sufficient enough to help our prediction.

2ND MODEL

The second model which I will be applying to the data set, in order to try and reduce my error value, will be an upgrade to my first model. I shall now examine how the users who decide what ratings to give to the movies will impact the RMSE value.

Compute the impact of all recorded users belonging to the edx data set. Compute penalty term b1; associated with users

```
user_rating <- edx %>%
  group_by(userId) %>%
  summarise(b1 = mean(rating - mu_cap))
user_rating_forecast <- mu_cap + validation %>%
  left_join(user_rating, by = 'userId') %>%
  .$b1
```

RMSE of Model 2

```
RMSE_MODEL_2 <- RMSE_Formula(a, user_rating_forecast)
RMSE_MODEL_2
```

```
## [1] 0.978336
```

INSIGHTS GAINED FROM MODEL 2:

I am pleased to see that the RMSE has decreased slightly after introducing the userId. This is because we are fed more information into our model, which helped to make a slightly more accurate prediction compared to Model 1.

3RD MODEL

After running the second model with the addition of users to our 1ST MODEL, we can see that the RMSE has decreased. This is what we want. For this 3RD MODEL, I shall observe the effect of including the 'movieId' to the model.

Compute the impact of all recorded "movieId's" belonging to the edx data set. Compute penalty term b2; associated with movieId's.

```
movie_user_rating <- edx %>%
  left_join(user_rating, by = 'userId') %>%
  group_by(movieId) %>%
  summarize(b2 = mean(rating-mu_cap-b1))
movie_user_rating_forecast <- validation %>%
  left_join(user_rating, by = 'userId') %>%
```

```

left_join(movie_user_rating, by = 'movieId') %>%
mutate(forecast = mu_cap + b1 + b2) %>%
pull(forecast)

```

RMSE OF MODEL 3

```

RMSE_MODEL_3 <- RMSE_Formula(a, movie_user_rating_forecast)
RMSE_MODEL_3

```

```
## [1] 0.8816096
```

INSIGHTS GAINED FROM MODEL 3:

Our error value has decreased compared to our previous model, but unfortunately not as much as I'd hoped for. I have also run out of variables to keep adding to my model. This means that I am going to have to turn to alternative methods in order to decrease my RMSE value to the desired result.

4TH MODEL

The third model, which includes ratings, userID & movieID, managed to obtain a lower RMSE reading compared to the first and second model, which is extremely reassuring that we're on the right path. I shall now implement the 'regularization' method to the data set in the hopes of obtaining an even lower RMSE reading than previously recorded.

Regularization

```

Regularization_Lamda <- seq(0, 50, 0.5)

Regularization_RMSE <- sapply(Regularization_Lamda, function(lam){
  Regularization_b1 <- edx %>%
    group_by(userID) %>%
    summarize(Regularization_b1 = sum(rating - mu_cap)/(n() + lam))

  Regularization_b2 <- edx %>%
    left_join(Regularization_b1, by = "userID") %>%
    group_by(movieId) %>%
    summarize(Regularization_b2 = sum(rating - Regularization_b1 - mu_cap)/(n() + lam))

  Regularization_b1_b2_mu_Forecast <- validation %>%
    left_join(Regularization_b1, by = "userID") %>%
    left_join(Regularization_b2, by = "movieId") %>%
    mutate(forecast2 = mu_cap + Regularization_b1 + Regularization_b2) %>%
    pull(forecast2)
  return(RMSE_Formula(a, Regularization_b1_b2_mu_Forecast))
})

```

RMSE OF MODEL 4

```

RMSE_MODEL_4 = min(Regularization_RMSE)
RMSE_MODEL_4

```

```
## [1] 0.8794441
```

INSIGHTS GAINED FROM MODEL 4:

On the bright side, the RMSE has decreased compared to our previous model. However, the error value is still not as low as I need it to be. This is extremely unfortunate as I was expecting to obtain the required

RMSE amount after implementing regularization. I will now have to apply another method now in the hopes of obtaining the required RMSE amount.

5TH MODEL

Seeing that our fourth Model (which used regularization) regrettably was unable to achieve and RMSE equal to or less than 0.8649; I have no choice but to try an alternative method to obtain this amount. I am now going to attempt the **Matrix Factorization** method in the hopes of obtaining the required RMSE value.

MATRIX FACTORIZATION

```
if(!require(recosystem)) install.packages("recosystem", repos = "http://cran.us.r-project.org")

## Loading required package: recosystem

library(recosystem)
set.seed(1)
matfac_recosys <- Reco()
valid_matfac <- with(validation, data_memory(rating = rating, user_index = userId,
                                              item_index = movieId))
edx_matfac <- with(edx, data_memory(rating = rating, user_index = userId, item_index = movieId))

matfac_parameters <- matfac_recosys$tune(edx_matfac, opts = list(niter = 30, nthread = 3,
                                                              dim = c(10, 30)))
matfac_recosys$train(edx_matfac, opts = c(matfac_parameters$min, niter = 30, nthread = 3))

## iter      tr_rmse      obj
##    0      0.9736  1.2040e+07
##    1      0.8726  9.9005e+06
##    2      0.8384  9.1694e+06
##    3      0.8166  8.7475e+06
##    4      0.8018  8.4744e+06
##    5      0.7903  8.2802e+06
##    6      0.7803  8.1266e+06
##    7      0.7721  8.0078e+06
##    8      0.7650  7.9102e+06
##    9      0.7589  7.8261e+06
##   10      0.7536  7.7576e+06
##   11      0.7489  7.6993e+06
##   12      0.7448  7.6484e+06
##   13      0.7410  7.6053e+06
##   14      0.7376  7.5664e+06
##   15      0.7344  7.5335e+06
##   16      0.7315  7.5000e+06
##   17      0.7288  7.4731e+06
##   18      0.7264  7.4469e+06
##   19      0.7241  7.4243e+06
##   20      0.7220  7.4054e+06
##   21      0.7200  7.3859e+06
##   22      0.7182  7.3669e+06
##   23      0.7164  7.3500e+06
##   24      0.7148  7.3352e+06
##   25      0.7134  7.3215e+06
##   26      0.7120  7.3092e+06
##   27      0.7106  7.2958e+06
```



```
## 28      0.7094  7.2861e+06
## 29      0.7081  7.2739e+06
```

```
matfac_solution <- matfac_recosys$predict(valid_matfac, out_memory())
```

RMSE OF MODEL 5

```
RMSE_MODEL_5 <- RMSE_Formula(a, matfac_solution)
RMSE_MODEL_5
```

```
## [1] 0.7808162
```

INSIGHTS GAINED FROM MODEL 5:

FINALLY! After running the data sets through Model 5, which uses the Matrix Factorization method, I have finally managed to obtain an RMSE value which is (equal to or) below 0.864999. The RMSE has decreased considerably compared to our previous Model 4. This method is able to generate a strong RMSE of 0.781.

RESULTS

I shall now create a 'results' table in which I shall record all of the RMSE Model recordings.

```
tab <- matrix(c(RMSE_MODEL_1, RMSE_MODEL_2, RMSE_MODEL_3, RMSE_MODEL_4, RMSE_MODEL_5), ncol=1,
              byrow = FALSE)
colnames(tab) <- c('RMSE Results')
rownames(tab) <- c('RMSE MODEL 1', 'RMSE MODEL 2', 'RMSE MODEL 3', 'RMSE MODEL 4', 'RMSE MODEL 5')
tab <- as.table(tab)
tab
```

```
##           RMSE Results
## RMSE MODEL 1    1.0612018
## RMSE MODEL 2    0.9783360
## RMSE MODEL 3    0.8816096
## RMSE MODEL 4    0.8794441
## RMSE MODEL 5    0.7808162
```

After completing 5 Model attempts with the sole purpose of obtaining an RMSE of less than 0.864999; I can happily say that I have managed to attain an RMSE result of 0.78065. Models 1-3 comprise of me adding more variables each time with hopes of decreasing the RMSE to the desired amount. Despite managing to get lower RMSE values each time as I progressed through these models, they still weren't low enough. I then had to turn to alternative methods such as regularization and ultimately to matrix factorization, which finally allowed me to obtain an RMSE below 0.864999.

CONCLUSION

This Capstone Project proved to be extremely joyful and insightful. It managed to test its students on numerous methods taught throughout this entire Data Science course. Moreover, it was able to encourage its students to go beyond what they've learned during this course and carry out additional research and learn new techniques and skills in order to complete this project. The chosen data set, MovieLens, was a fun and relatable topic which kept its students engaged and interested (it definitely convinced me to watch a few movies which I've never seen before).